



Trinity College Dublin

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

School of Engineering

Electronic Engineering

MERIT.jl: Julia's Version

Aaron Dinesh

Supervisor: Associate Prof. Declan O'Loughlin

April 7, 2024

A Final Year Project submitted in partial fulfilment
of the requirements for the degree of
MAI (Electronic and Computer Engineering)

Declaration

I hereby declare that this Final Year Project is entirely my own work and that it has not been submitted as an exercise for a degree at this or any other university.

I have read and understand the plagiarism provisions in the General Regulations of the University Calendar for the current year, found at <http://www.tcd.ie/calendar>.

I have completed the Online Tutorial on avoiding plagiarism 'Ready Steady Write', located at <http://tcd-ie.libguides.com/plagiarism/ready-steady-write>.

I consent/do not consent to the examiner retaining a copy of the thesis beyond the examining period, should they so wish (EU GDPR May 2018).

I agree that this thesis will not be publicly available, but will be available to TCD staff and students in the University's open access institutional repository on the Trinity domain only, subject to Irish Copyright Legislation and Trinity College Library conditions of use and acknowledgment. **Please consult with your supervisor on this last item before agreeing, and delete if you do not consent**

Signed: Aaron Dinesh Date: April 7, 2024

Abstract

MERIT aims to provide a software framework that is robust, easy to use and performant. It implements a variety of microwave imaging algorithms and a myriad of helper functions, all while leveraging the powerful features available in Julia. MERIT.jl also implements a “Scan” abstract datatype which allows users to subtype their own specialized datatype. Organizing the datatypes in this way means that MERIT.jl plays very well with Julia’s own type hierarchy and also the other language features that depend on this. To encourage type safety, MERIT.jl implements a lightweight Points class which allows for efficient processing of coordinate points. In this way, collections of points won’t simply be a matrix of Floats or Ints instead, they would be a Vector of the Points type. In this way, the Julia compiler will throw an error when Points aren’t passed in the right argument, instead of providing a wrong output.

Acknowledgements

Thanks, Everyone!

List of Figures

1.1	Example of a Fully Multistatic Configuration (Top-Down)	3
1.2	The MARIA M4 and M5 system. (a) The MARIA M4 antenna array. (b) The M4 in a clinical setting. (c) The integrated M5 package [2]	3
1.3	The Wavelia System [3]	4
1.4	The Leveled Multistatic Approach of Wavelia [3]	5
1.5	The TSAR Prototype [7]	6
1.6	Phased Array Diagram	7
1.7	Point Emitter with a Phased Array	8
1.8	Comparason between DAS (a, c) and DMAS (b, d) Beamformers	10
2.9	Multiple Dispatch (top) vs Single Dispatch (bottom)	15
2.10	Type hierarchy for the Integer Type	16
2.11	Current type hierarchy in MERIT.jl	18

Contents

Introduction	1
Background	2
1.1 Literature Review	2
1.1.1 MARIA M4	2
1.1.2 Wavelia	4
1.1.3 TSAR	5
1.1.4 Beamformers	7
Delay and Sum	7
Weighted Delay and Sum Beamformer	9
Delay Multiply and Sum	9
Julia	11
2.1 Why Julia?	11
2.1.1 Python	11
2.1.2 MATLAB	12
2.1.3 Julia	12
2.2 Features in Julia	13
2.2.1 Multiple Dispatch	13
2.2.2 Type Heirarchy	15
2.2.3 Parametric Polymorphism	18
2.2.4 Type Stability	19

Introduction

MERIT.jl was motivated by a need to increase the interoperability between the data produced from systems and the programs that would be able to process them. Currently, any researcher who wants to design a novel microwave imaging algorithm or test the efficacy of their new systems has to spend time and effort creating entirely new software suites. More often than not these programs either tend to be closed source, or not compatible with other systems mainly because there are not many incentives for a researcher to invest the time required in creating a truly configurable and extensible open-source system. With the rise and success of clinical trials in microwave imaging, this problem is only going to get worse as more and more systems are developed.

The use of such open-source software has seen widespread use in a myriad of fields. PsychoPy and PsyToolkit are two such frameworks that revolutionized the field of psychological sciences. By implementing commonly used functions and scripts, they have allowed researchers to design and run experiments in a matter of hours. It has also allowed researchers who have little to no programming experience to get up and running with automated data processing, thereby allowing them to focus more on the quality of their experiment [1].

That is where this project comes in. MERIT.jl aims to be an easy-to-use, extensible and featureful library. The goal being that anyone, regardless of coding experience, would be able to quickly create a script that allows them to process and visualize the scan data they have collected, as well as allowing more experienced coders to develop and test new algorithms with ease. The following sections will contain:

- A literature review on the existing microwave imaging systems
- A look into existing microwave imaging frameworks
- A discussion about the design choices and Julia features that are included
- An examination of the results and possible future work

The source code and the library are hosted on GitHub for anyone to view and amend: <https://github.com/AaronDinesh/MERIT.jl>

Background

1.1 Literature Review

In order to understand the design requirements of the library, a literature review needed to be performed. Due to the relative infancy of the field, there are various competing configurations with no clear winner in sight. However, some patterns can still be gleaned from current clinical trials which are broadly representative of the direction the field is moving in. In all of these systems, the patient to be scanned would lie in the prone position on an imaging table. The patient would then pass their breast through a hole in the table into some type of imaging apparatus. The imaging apparatus varies slightly from system to system, but they all contain some type of antenna array that is used for the imaging process. Most systems will adopt one of the following antenna setups

1. Monostatic
2. Leveled Multistatic
3. Fully Multistatic

These terms relate to the position of antennas around the breast tissue and how the resulting scan data would be structured.

1.1.1 MARIA M4

The first system we will look at is the MARIA M4 system developed by Preece et al within the Electrical and Engineering Department of the University of Bristol [2]. This is the 4th iteration in a series of MARIA systems that evolved from a system of 16 UWB antennas to 60 antennas in the current system. These all operate in a multistatic configuration, meaning that any antenna in the array can listen to any other antenna in the array, an example of which can be seen in Figure 1.1. This figure shows a top-down view, however, one can imagine this being generalized to a hemisphere of antennas around the breast.

As stated before, the MARIA M4 system makes use of the UWB spectrum over a frequency range of 3.0 to 8.0 GHz. A commercial-grade Vector Network Analyzer (VNA) was used as the system signal source. The VNA, operating in a stepped continuous wave mode, was used to both produce the transmitted sine waves and record the frequency and phase of the reflected waves at the receiving antennas. The choice of a commercial-grade VNA is indicative of the prototype nature of the MARIA M4 system, where easy reconfigurability of the frequency imaging range would be of higher importance than price. While not explicitly stated, one would presume that the M5 system would replace the VNA with some application specific




Figure 1.1: Example of a Fully Multistatic Configuration (Top-Down)




Figure 1.2: The MARIA M4 and M5 system. (a) The MARIA M4 antenna array. (b) The M4 in a clinical setting. (c) The integrated M5 package [2]

hardware to reduce the overall cost of the system. The M4 system exploits the inherent symmetry in the antenna reciprocity to halve the number of channels (made of a transmitting and receiving antenna) collected, thereby speeding up the scan time. For the MARIA M4 system, this equates to a 1770 reduction in the number of channels collected. Figure 1.2, shows the antenna array used in the M4 system (a), as well as the M5 system (c) which is an integrated package.

The team also conducted a clinical study in order to test the efficacy in women who already attend symptomatic breast care clinics. In total 86 patients were included with ages ranging from 24 - 78 years old; the mean age being 51.4. The inclusion criteria for the study required that possible participants be clinically symptomatic, be able to be imaged by ultrasound and mammograms (these being the control), be able to lie in the prone position and have cup sizes within 310 to 850 ml. The types of lesions included in the study were mostly cysts and cancers, but some other conditions such as hematoma, lipoma and fibroadenoma were also included. The goal of the study was to test the sensitivity of the M4 system; the sensitivity metric being determined based on the ability of the system to localize a lesion as it correlated with the location in the ultrasound image. The M4 system showed a sensitivity of 74% (64/86) when compared with the "gold-standard" of an ultrasound. The research team also divided the group into pre-/peri- and post- menopausal women and found that the sensitivity was 75% and 73% respectively. However, the reliability of these results are called into question when considering the sample size of the study. Given a sample size of 86, assuming a normal distribution and that the results are statistically significant ($p < 0.05$, $Z = 1.96$), a 7.11% margin of error was calculated. While this may not be enough to conclusively prove that the



Figure 1.3: The Wavelia System [3]

M4 system is a viable alternative to Mammograms, it is enough to show promise. With a larger sample size, this margin of error could be narrowed further.

1.1.2 Wavelia

The second system that was considered was the Wavelia Microwave Breast Imaging System developed by MVG Industries [3]. The Wavelia system integrates the imaging system as well as the examination bed into one complete package (Figure 1.3). The integrated package makes the Wavelia system an appealing choice for some hospitals, however, its large size may be a barrier to adoption in some facilities where space is a premium.

Like in the MARIA M4 system, patients lie in a prone position and place their breasts in the circular cutouts on the bed. The system then begins to create a 3D reconstruction of the exterior of the breast using a stereoscopic camera. This also allows for the breast volume to be explicitly calculated rather than being inferred like in the MARIA system. The Wavelia makes use of the UWB spectrum when imaging the breast, although they opt for a narrower frequency range of 0.5 - 4.0 GHz compared to the 3.0 - 8.0 GHz range of the MARIA system. The antenna configuration, unlike the MARIA system, is an array of 18 Vivaldi-type probes arranged in a concyclic manner on a horizontal plane. These antennas operate in a Multistatic manner and image the breast in sections parallel to the coronal plane. The entire antenna assembly moves downwards in 5mm intervals to image the entire breast (Figure 1.4). This is a leveled multistatic system as opposed to the fully multistatic system in the MARIA M4. This leveled multistatic approach has the benefit of a theoretically infinite vertical resolution. If the radiologist wanted a finer resolution along the coronal plane, they would just have to tweak the array vertical step size, rather than having to manufacture an entirely new antenna array, like in the MERIT system. This leveled approach also allows the parameters of the reconstruction algorithm to be easily changed based on the particular section of the breast that we are imaging. The fully multistatic approach, however, would need significant additional logic in the post-processing steps to determine which channels are coplanar with a particular section.




Figure 1.4: The Levelled Multistatic Approach of Wavelia [3]

The Wavelia paper [3] conducted a feasibility study on 25 female participants who were recruited after presenting with symptoms to the Symptomatic Breast Unit at Galway University Hospital, Ireland. Their inclusion criteria required that participants to have a mammogram that was performed at the time of symptomatic presentation to the hospital, be capable of lying in the prone position for a length of 15 minutes, have a bra size larger than 32B and a breast size equivalent to a B cup. The final criteria required that the participant's submerged breast have enough margin between the breast and the cylindrical container in order to accommodate the transition liquid. This was determined by the clinician at the time of the trial. Out of the 25 patients who presented with a palpable lump, one patient's scan had to be removed, due to the lump being later classified as normal tissue. 11 of the participants had a biopsy confirmed carcinoma and out of these, the Wavelia system detected 9 lesions, with 7 being located to the appropriate region. Overall the system detected an abnormality in 21 of the 24 participants, leading to a sensitivity of 87.5%. The researchers do note some limitations of the Wavelia system, namely it can't detect any lesions smaller than 10mm. This is significant since the size of the detected lesion plays a big factor when deciding whether a lesion is cancerous or not. Another limitation of the system is that breast sizes that are too small cannot be scanned in any great detail by the Wavelia system. Due to the patients being in the prone position, their breast tissue needs to hang down far enough to have multiple sections of their breast be imaged by the antenna array. The researchers are working on a subsequent system that should address all the aforementioned limitations. Overall the participants had a positive outlook on the system. 23 out of the 25 women said that they would recommend the procedure to other women and all of the women agreed that the information provided was clear and well understood.

1.1.3 TSAR

The third and final system considered for the project was the TSAR system [4]. Standing for Tissue Sensing Adaptive Radar, this system was developed by the University of Calgary to address some of the shortcomings of the previous two systems. In the MERIT and Wavelia systems, reflections from the skin can muddy the data and can lead to artifacts in the final image. In order to get around this, both systems take an additional scan, offset by a fixed rotational amount in the coronal plane. The idea being, that any reflections that appear on the first scan would appear with similar amplitude in the second scan at the same time position,



Figure 1.5: The TSAR Prototype [7]

while the tumor would show in a different place, provided that the tumor doesn't lie on the axis of rotation. This method is known as Rotational Subtraction [5]. The TSAR system, on the other hand, makes use of an adaptive algorithm that instead estimates the skin response at an antenna as the weighted sum of the responses from the neighboring antennas. This skin response can then be subtracted from the current antenna to remove the reflection artifacts [6].

Like the previous two systems, patients lie in the prone position on the examination table, with their breast being submerged in an immersion liquid. However, unlike the previous two systems, the TSAR operates in a monostatic configuration. In this setup, there is only one antenna that acts as both the transmitting and the receiving antenna. This antenna, usually fixed on some type of rotating apparatus, would rotate around the breast to image a section. It would then step vertically by a fixed amount and repeat the previous step, eventually imaging the entire breast. The TSAR system also operates on a much wider section of the UWB spectrum, from 50 MHz to 15 GHz, with the frequency data being collected by a VNA. The system also makes use of a laser in order to help with the 3D reconstruction of the breast during post-processing. The prototype setup can be seen in Figure 1.5

The benefit of the monostatic configuration is that the number of antenna locations per row and the number of rows are parameters that can be tweaked depending on whether the radiologist would want a quicker scan or higher precision. One can imagine the scenario of a rapid screening center as part of a national screening program where the speed of the scan is valued over the precision of the scan. As such this provides a notable benefit over the previous two systems. The clinical trial conducted by this study was extremely limited, only including 8 successfully imaged patients. Due to the low sample size, one can't reliably say whether this system and its adaptive reflection-suppressing algorithm can outperform the other two systems in terms of accuracy, although the TSAR system does show promise.

1.1.4 Beamformers

Without loss of generality, imagine a line array of wave emitters. If all of these begin to emit at the same time, the individual waves would constructively and destructively interfere with each other as their peaks and troughs align and misalign. This property can be exploited in such a way so that the waves constructively interfere in one direction and destructively interfere in all the other directions, effectively aiming the beam in a particular direction. A diagram of this can be seen in Figure 1.6.




Figure 1.6: Phased Array Diagram

This is the forward beamforming process, but one can also consider the reverse process. Considering Figure 1.7 imagine a point emitter embedded in a 2D plane, that radiates circular waves evenly in all directions. Now imagine a line array that receives this wave. Due to the diffusion of the wave through space, the same wave will appear at each antenna at different times. This would manifest as the same amplitude appearing at different times, in the signal graphs, even though the impulses come from the same wave. Reverse beamforming then, is the process of varying the phase and amplitude of the received signals in order to estimate the intensity at the location where the wave originated. Inverse beamforming is a process used in many fields such as radio astronomy and seismic imaging, so there already is a wealth of research and plenty of algorithms to choose from. The rest of this section will talk about the various beamformers that are popular in the field of microwave imaging.

Delay and Sum

The Delay and Sum (DAS) beamformer is the template for most other beamforming algorithms in the field. The DAS beamformer posits that every antenna has recorded the same source and that the delay in each signal is due to the relative distance between the receiving antenna (\mathcal{A}') and the transmitting antenna (\mathcal{A}). As such, if one was to delay the signals by the correct amount, and sum over all the received signals, one would be able to estimate the energy at the source. This same idea can be applied to signals received from the aforementioned imaging systems. Since these work in the frequency domain, the following explanations and equations




Figure 1.7: Point Emitter with a Phased Array

will work on this assumption, but these equations can easily be converted to the relative time domain functions. As stated before a point from the imaging domain is chosen (r). We then estimate the path delay of the wave from the transmitting antenna to the receiving antenna along the path.

$$\tau_{\mathcal{A}'_i, \mathcal{A}_j}(r) = \frac{\epsilon}{c_0} [\|\mathcal{A}' - r\| + \|r - \mathcal{A}\|] \quad (1.1)$$

ϵ is the relative permittivity of the medium. In reality, this value changes from patient to patient and even varies within the breast of each patient depending on the path taken. However in order to achieve a practical beamformer, one must estimate an averaged relative permittivity value for the entire breast, we will label this as ϵ_i and it will parametrize the DAS beamformer. As well as this, equation 1.1 assumes a straight-line path between the point and the transmitting and receiving antennas. This is rarely the case, however, in practice making this assumption only yields a maximal error of 3mm in position while greatly simplifying the delay calculations [?]. Using this, the received signals are delayed and then summed across all stepped input frequencies. This result is then squared to give us the energy at the chosen point. Thus the point is imaged, This process is repeated for all the points in the imaging domain, overall the DAS beamforming equation can be represented as such:

$$I_{\epsilon_i}(r) = \left[\sum_{\Omega} \sum_{\mathcal{A}'} \sum_{\mathcal{A}} S_{a,a'}[\omega] e^{j\omega \tau_{\epsilon_i, a, a'}(r)} \right]^2 \quad (1.2)$$

The DAS algorithm was the chosen beamformer for the MERIT.jl library due to its simplicity to implement and the time constraints I had to work with. However, two other beamformers will be discussed below and will be compared with the DAS algorithm.

Weighted Delay and Sum Beamformer

The Weighted Delay and Sum (WDAS) Beamformer can be considered as a further generalization of the DAS beamformer. In equation 1.2 all channels contribute equally to the final result due to the implicit weighting factor of 1. The equation also assumes that all signals travel along similar paths, and therefore a constant speed (due to the fixed ϵ_i). In reality, this is only true for antennas that are closer together. The greater the distance between the transmitting and receiving antennas, the greater the chance, that the waves deviate from the straight line path, ergo the estimation of the speed and subsequently the delay along that path will be wrong. Essentially, the algorithm needs to have some mechanism that would penalize the signal from antennas that are further away and reward signals from antennas that are close by. This is what S.A. Shah Karam et al proposed in their 2021 paper, "Weighted delay-and-sum beamformer for breast cancer detection using microwave imaging" [?]. They define a weighting factor based on the transmitter-receiver distance (TRD_i) for the t^{th} observation:

$$w_i = \alpha - |r_{T_{r_i}} - r_{R_{c_i}}| \quad (1.3)$$

In order to reward signals from antennas that are close and penalize antennas that are far away one must subtract the TRD_i from a positive constant, α . This parameter is patient-specific and must be changed based on the homo- or heterogeneity of the breast tissue. The paper used an α of 20cm for their results. One important thing to note is that these weighting factors are data-independent. Since the TRD_i relies only on the distance between the antennas, these weighting factors can be computed beforehand. Another thing of note is that these weighting factors are independent of the chosen focal point as well, which means we can compute a normalized weight and apply it to the signal before employing the traditional DAS algorithm. The normalized weighting factor \hat{w}_i is calculated as follows, where M is the number of channels:

$$\hat{w}_i = \frac{w_i}{\sum_{i=0}^M w_i} \quad (1.4)$$

Delay Multiply and Sum

The Delay Multiply and Sum (DMAS) algorithm was first proposed by Hooi Been Lim et al in 2008 [?]. Instead of considering each channel in isolation across the frequency bands, the DMAS algorithm proposes a pair-wise multiplication of the delayed signals before summing across the frequency range. For a focal point r the equation would be as follows:

$$I_{\epsilon_i}(r) = \sum_{\Omega} \sum_{\mathcal{A}'} \sum_{\mathcal{A}} \sum_{\mathcal{B}'} \sum_{\mathcal{B}} S_{a,a'}[\omega] e^{j\omega\tau_{\epsilon_i,a,a'}(r)} S_{b,b'}[\omega] e^{j\omega\tau_{\epsilon_i,b,b'}(r)} \quad (1.5)$$

Here $\mathcal{A}' = \mathcal{B}'$ are the receiving antennas and $\mathcal{A} = \mathcal{B}$ are the transmitting antennas. The paper then compared the DMAS algorithm with the traditional DAS algorithm and found that the pairwise multiplication of the delayed signals before summation provided greater contrast and sharper results, greatly increasing the SCR of the image. Figure 1.8 shows the differences between the DAS and DMAS algorithms and most importantly, how the DMAS beamformer provides a greatly reduced noise floor.




Figure 1.8: Comparason between DAS (a, c) and DMAS (b, d) Beamformers

The original paper never provided an explanation for the surprising effectiveness of the algorithm, leaving the explanation for a future paper that they never wrote. However, this did not stop the wider research community from providing possible explanations as to how the DMAS algorithm provides the results it does. One explanation provided by Prof O'Loughlin posits that the pairwise multiplication rewards signals that have a high degree of coherency while disproportionately penalizing incoherent signals. In essence, the coherent signals get brightened a lot more than the incoherent signals would [?]. G. Matrone et al backs up this idea by noting that the DMAS beamformer, after the signals are time aligned like in DAS, is the autocorrelation function of the receiver antenna with the auto terms excluded [?].

Julia

2.1 Why Julia?

One of the goals of MERIT.jl was to streamline the development process for new imaging algorithms. As well as this MERIT.jl aims to lower the barrier of entry and the coding knowledge needed in order to set up these data processing and visualizing pipelines. In order to fulfill these dreams, the right coding language needed to be selected. Since the barrier of entry needs to be low, a high-level language needs to be selected and since the primary target of MERIT.jl are researchers, it needs to be a language that they are familiar with, or one which they can easily learn. If the language chosen is too difficult to learn, any potential time savings from the ease of use would be heavily outweighed by the time spent learning the language.

2.1.1 Python

Python was one of the languages considered. Python's philosophy was code readability over all and this fact is complimented by its high-level syntax and indentation requirements. It was first introduced towards the end of the 1980's by Guido van Rossum and has had an expansive community ever since. The Python package repository and pip have allowed the community to develop libraries that can be easily shared and downloaded. All these factors catapulted Python into the limelight and it has become the go-to language for any data science and deep learning application. However, several drawbacks severely limit the usability of Python for high performance applications. Firstly, Python is an interpreted language, meaning that at runtime, the Python interpreter reads the Python file line by line and calls the relevant machine code. Due to this interpretation step, raw Python code is slow to run. Due to this reason, most performant libraries in Python tend to use C-optimized implementations behind the scenes. So in order to develop any decently performing Python libraries, one also has to be proficient in C as well as know how to create C and Python bindings. Secondly, the Python Interpreter makes use of a runtime lock known as the Global Interpreter Lock (GIL) which makes parallel processing in Python incredibly difficult. The reason for the GIL comes down to Python's use of reference counting for memory management. Reference counting is a method where each object in Python gets assigned a reference variable that keeps track of the number of references that point to that object. As references to the object get created, the count goes up, as references get deleted or reach the end of their scope, the count goes down. When the reference variable reaches zero, this becomes a flag for the object to be deleted, since there are no longer any references pointing to this object. In some threaded applications, this can cause some issues where the reference count gets updated by multiple threads simultaneously. In some cases, this can cause the reference count to never reach zero

leading to a memory leak, or it can reach zero too early and the object gets removed. To get around this, Python puts a lock on the interpreter itself, this is the GIL. This means that any Python bytecode needs to first acquire a lock on the interpreter before it can be executed. This makes multithreading in Python incredibly difficult and slower than it would be in other languages. So for these reasons, Python was rejected as the software of choice.

2.1.2 MATLAB

MATLAB was another language that was considered. Developed in 1984, it became the goto software for many research purposes due to its ease of use and intuitive operations on matrices and its multi-dimensional analogues. It is taught across every single engineering college, and one company estimates that MATLAB is being used in over 57,811 companies [?]. MATLAB, however, is still an interpreted language and therefore can be quite slow some times. A study conducted by Aruoba and Fernández-Villaverde found that their MATLAB code ran about 3x slower than the same code written in C++, highlighting just how big of a difference a compiler can make [?]. It should be noted that the MATLAB version of the code ran about 30.26x faster than the same Python code, implying that even though both are an interpreted language, the MATLAB interpreter is much more optimized than the Python interpreter. But one of the biggest drawbacks by far, is the fact that MATLAB is a license based language. In order to use MATLAB, one must pay a yearly subscription of anywhere from €120 to €3,650 depending on the purpose for which it is used. This goes directly against the open-source vision of MERIT.jl. This cost can become prohibitively expensive for small teams trying to work on new systems and algorithms. Octave was briefly considered as it is a free and open-source competitor to MATLAB, but this idea was quickly dropped when it became clear that Octave's main goal was compatibility with MATLAB scripts rather than performance over MATLAB. For these reasons, MATLAB was also rejected as the software of choice.

2.1.3 Julia

The third and final language considered was the Julia programming language. Julia was developed by Jeff Bezanson, Stefan Karpinski and Viral B. Shah, and was first released to the public in 2012. Julia was designed to be a coding language that offered high-level syntax while also automatically compiling the code at runtime. Julia's Just In Time (JIT) compiler, allows it to offer the full dynamism of Python and MATLAB while avoiding the drawbacks that come with having to use an interpreter. The aforementioned study found that Julia was only 1.47x slower than the comparable C++ implementation. Julia also offers features that are not in Python or MATLAB such as parametric polymorphism, multiple dispatch, efficient garbage collection and a JIT that is capable of optimizations and ahead-of-time compilation. Julia also bakes in native support for multiple parallel programming paradigms such as GPU programming and multithreading, making it an appealing choice for people working on high-performance compute clusters. One other feature that made Julia wildly popular was the native ability to call C and Fortran libraries without having to create any special wrappers. This "It Just Works" ideology is what garnered Julia its well deserved popularity. With its rising fame and easy to understand syntax, this was the coding language decided for MERIT.jl. In the next few sections, the various Julia features used will be discussed in more detail.

2.2 Features in Julia

2.2.1 Multiple Dispatch

Multiple Dispatch is a programming paradigm in which a function can be dispatched (invoked) based not only on the function name but also on the type and order of input arguments. This is opposed to the single dispatch programming paradigm where the function that is dispatched usually depends on a special argument before the function name. In almost all programming languages this is the variable name for that class. For example, consider the following code:

```
1      Class Dog:
2          name::string
3
4          function says(a::string):
5              print("The dog, $self.name says $a")
6          end
7      end
8
9      Class Cat:
10         name::string
11
12         function says(a::string):
13             print("The cat, $self.name says $a")
14         end
15     end
16
17     billy = Dog("billy")
18     kate = Cat("Kate")
```

If we wanted to call the "says" function for the Dog or Cat class, then we would have to write, `billy.says("Hello World")` or `kate.says("Hello World")`. This is the single dispatch paradigm and it is probably the one most people are familiar with. This fits well into an object-oriented programming language, where objects are used to encapsulate concepts and ideas and therefore we create specialized functions that operate on the data contained within that class. However, one drawback is that the compiler relies on the user to remember which methods belong to which class, and also which methods are callable. Also, it isn't clear how one would write a function that allows for the Cat and the Dog to interact with each other. Instead, now consider the multiple dispatch paradigm. In Julia particularly, it is important to note that methods no longer belong to objects as we were forced to do in single dispatch. Added to this, methods no longer have to be defined within the "class" but can be done after class declaration or even class instantiation. Consider the example below:

```
1      struct Dog{
2          name::string
3      }
4
5      struct Cat{
6          name::string
7      }
```

```

8
9     function says(pet::Cat, a::string)
10         print("The cat, $pet.name says $a")
11     end
12
13     function says(pet::Dog, a::string)
14         print("The dog, $pet.name says $a")
15     end
16
17     function encounter(petA, petB)
18         print("$petA.name encounters $petB.name and $meets(
19 petA, petB)")
20     end
21
22     function meets(petA::Cat, petB::Dog)
23         return "hisses"
24     end
25
26     function meets(petA::Dog, petB::Cat)
27         return "barks"
28     end
29
30     billy = Dog("Billy")
31     kate = Cat("Kate")
32     says(billy, "Hello")
33     says(kate, "Hello")
34     encounters(kate, billy)
35     encounters(billy, kate)

```

Here we have a similar situation as before, but instead of having classes, we have structs and methods that take those structs as arguments. On lines 30 and 31 in the code block above, we get a similar output as in single dispatch. But where multiple dispatch really shines is in the next two lines. Here we have the cat and the dog encountering each other. In single dispatch, someone would have to create a wrapper library to handle the interaction between two classes, but in multiple dispatch it is as easy as creating another method. This can be seen in the "meets" functions defined in the code block above. Since multiple dispatch relies on the input arguments as well, we can create specialized "meet" functions for when a Dog meets a Cat or when a Cat meets a Dog. In the above code block executing line 32 would print out "Kate encounters Billy and hisses". This is because we had a struct of type Cat as the first input argument to "meets" and a struct of type Dog as the second argument. Line 33 on the other hand, would print "Billy meets Kate and barks", since the order of the input arguments is reversed. Another notable consequence from this is that if anyone else wanted to extend "meets" and therefore "encounters" to other animals, they would only need to define a struct with the same fields as Cat and Dog, and new meets methods as such:

```

1     struct Rabbit{
2         name::string
3     }
4

```

```

5  function meets(petA::Cat, petB::Rabbit)
6      return "chases"
7  end
8  function meets(petB::Rabbit, petB::Cat)
9      return "runs away"
10 end
11
12 kate = Cat("Kate")
13 johnny = Rabbit(Johnny)
14 encounters(kate, johnny)
15 encounter(johnny, kate)

```

Line 12 in the above code block would print "Kate encounters Johnny and chases" while line 15 would print "Johnny encounters Kate and runs away". While these examples may be simple, they showcase the powerful flexibility and extensibility behind multiple dispatch. It is this feature that allows all the Julia libraries to "just work" together, without needing any glue-code. Some other developer could come along and create a new package for a "Lion" and so long as it implements a struct with the same required field names as the other animals and creates a "meets" function, any variable of type Lion will work well with the "encounters" function.

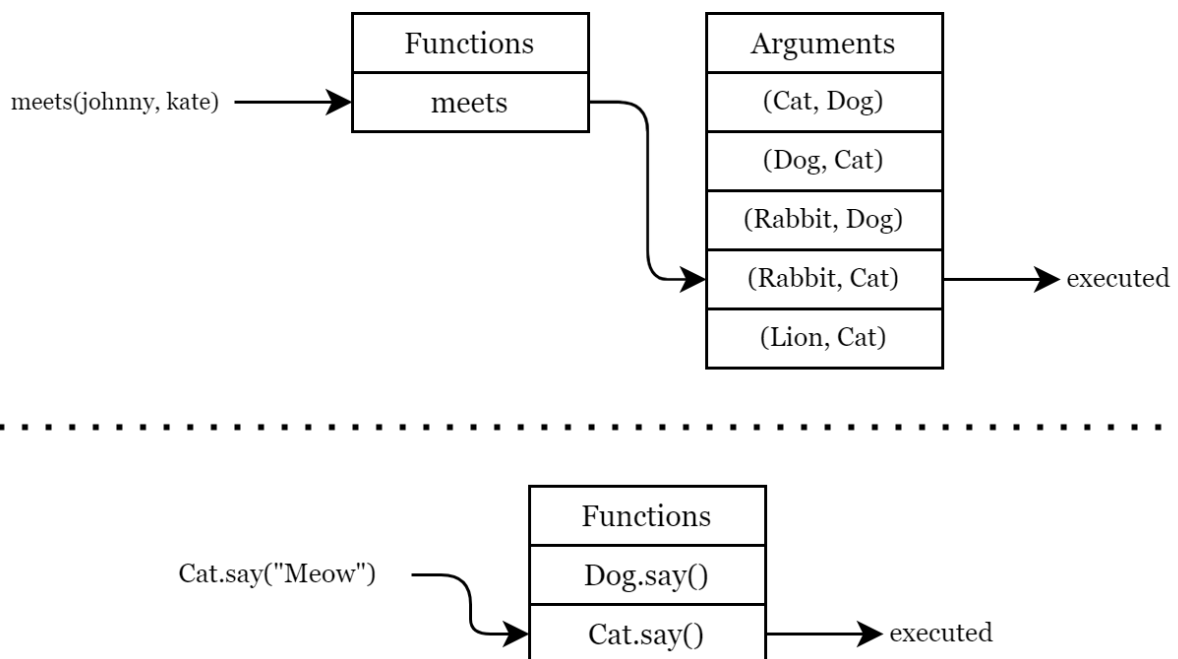


Figure 2.9: Multiple Dispatch (top) vs Single Dispatch (bottom)

2.2.2 Type Hierarchy

All the types in Julia are arranged in a tree-like structure and can be broadly classified into two categories, an Abstract Type or a Concrete Type. Abstract types are the internal nodes of the type tree, having both parents and children, while concrete types are the "leaves" of this tree. One notable difference between abstract types and concrete types is that abstract types cannot be instantiated; they serve only as nodes in the type graph. While this may seem




Figure 2.10: Type hierarchy for the Integer Type

pointless at first glance, it is this very feature that allows the multiple dispatch paradigm mentioned before to really shine. Shown in Figure 2.10, is the type hierarchy for the Integer type, the abstract types are highlighted in red, whereas the concrete types are highlighted in blue.

As well as giving us a way to logically organize types, the type hierarchy tightly integrates with the multiple dispatch system. Consider for example the code block below:

```

1  abstract type Animal end
2
3  struct Cat <: Animal
4      name::string
5  end
6
7  struct Dog <: Animal
8      name::string
9  end
10
11 function encounters(petA::Animal, petB::Animal)
12     verb = meets(petA, petB)
13     println("$(petA.name) meets $(petB.name) and $(verb)"
14 )
15 end
16
17 meets(petA::Animal, petB::Animal) = "passes by."
18 meets(petA::Cat, petB::Dog) = "hisses"
19 meets(petA::Dog, petB::Cat) = "barks"

```

```

19     whiskers = Cat("Whiskers")
20     fudge = Dog("Fudge")
21     encounters(whiskers, fudge)
22
23
24     #####
25     #Defined in another library that subtypes the Animal
26     abstract class.
27     struct Rabbit <: Animal
28         name::String
29     end
30
31     chomper = Rabbit("Chomper")
32     encounters(whiskers, chomper)

```

In the above code block, we define an abstract `Animal` type, by default this is a subtype of the abstract type `Any`. Then we define a concrete type `Cat` and a concrete type `Dog`. When line 22 gets executed, the multiple dispatch system will choose the correct `meets()` function and "Whiskers meets Fudge and hisses", this is as expected. Now another developer might want to implement a `Rabbit` type. So they create the `Rabbit` type as shown above, and specifies that it is a subtype of the `Animal` abstract type via the `<:` operator. Even if they do not implement a `meets()` that is specialized for their `Rabbit` class, the `encounters()` will still work since it only expects an input that is a subtype of `Animal`. The `meets()` function will also work, since the original author implemented a `meets()` method that accepts arguments of type `Animal` or any subtypes of it. But why didn't the earlier call to `encounters()` (and consequently `meets()`) fail, since both lines 15 and 16 are valid options for the dispatcher? This is an example of where multiple dispatch benefits from having a type hierarchy. The Julia dispatcher always dispatches the function that is most specific across **ALL** its arguments. So in the `Cat` and `Dog` case, the `meets()` on line 16 will be dispatched. In the `Cat` and `Rabbit` case, the `meets()` on line 15 is the most specific function and so that gets dispatched. This allows for amazing extensibility that isn't found in many other languages. Often it is impractical for a developer to think of all the ways their library would be used or all the types or functions that the end-users would require. In Julia, they don't have to think about it. Instead, they can define an abstract class and some methods that accept the abstract class as inputs. The users can then create subtypes of this abstract class and define their own methods for their new type. Neither developer has to be worried about whether they were thorough enough to consider all possibilities, they can write code for their own purposes and trust that the Julia dispatcher and type hierarchy would automatically determine the correct function to call.

This idea is heavily used in the `MERIT.jl` library. Currently, most of the research in microwave imaging is centered around the breast and breast cancer, but in the future, this could be implemented for other body parts. As such, the library needs to be flexible to allow for easy updates. `MERIT.jl` is centered around the `Scan` abstract type, from this type we subtype the `BreastScan` type, which holds all the information from the scan of a particular breast. If in the future microwave imaging gets used for a chest scan, all a developer would need to do is to implement a `ChestScan` struct, subtyped from `Scan`, and implement a few functions and the Julia dispatcher will handle the rest. I, as the original developer, do not need to worry about the intricacies of how a `ChestScan` struct would work, or what fields it would need. I can trust

that future developers can easily extend this library without introducing breaking changes to the library core. In this way, MERIT.jl achieves extensive flexibility and expandability which would not be possible in other languages.




Figure 2.11: Current type hierarchy in MERIT.jl

2.2.3 Parametric Polymorphism

Parametric polymorphism refers to the programming paradigm where a function can be created generic to the input argument type. The compiler can then strongly type the function with the correct types. Consider the simple example of an `addtwo(x)` function.

```

1  function addTwo(x::T) where {T <: Real}
2      return x + T(2)
3  end

```

Instead of restricting `x` to a specific type, we replace it with the variable `T` and instead place a restriction on `T`. In this case we say that `T`, and by extension `x`, can be any type that is a subtype of the abstract `Real` type. We also convert the `2` from an `Int64` type, to whatever type `T` is, to allow for accurate summation. If parametric polymorphism was not a feature in Julia, this function would need to be duplicated 16 different times to account for all the concrete subtypes of `Real` (see Figure 2.10). Parametric Polymorphism with the Julia type hierarchy allows a developer to have incredible expressive power with very few lines of code. MERIT.jl uses this feature extensively due to the fact that the library needs to be agnostic to the type of the input data. When instantiating the `BreastScan` struct the user can provide the types of the data that they will load and can therefore set the internal data type of the struct. It would be impossible to support all 4,176 different permutations of every single function, however, with parametric polymorphism, the function only needs to be defined once and the Julia compiler will handle the rest. This means that researchers and developers don't have to worry about whether the library functions can support the data type of their data, so long as it follows the type restriction set on the function, it will produce an output, greatly simplifying the coding experience.

2.2.4 Type Stability

Type Stability is a coding discipline that Julia recommends all code to use. A function whose contained variables have a consistent type for the lifetime of that function is considered to be "type stable" or to have "type stability". Another way to think about this is that if the output type only depends on the input types, then the function is type stable. Consider the following two functions:

```
1  function unstable()
2      x = 1
3      for i = 1:10
4          x /= rand()
5      end
6      return x
7  end
8
9  function stable()
10     x = 1.0
11     for i = 1:10
12         x /= rand()
13     end
14     return x
15 end
```

In both functions, we have a variable `x` which gets divided by a random float 10 times, with the result being returned. However, the first function is considered to be type unstable while the second function is type stable. This is because in the first function `x` was declared as an `Int64`, but when we divide by the random `Float64`, the data type of `x` changes to `Float64` as Julia performs floating point division by default. Whereas in the second function, `x` gets declared as a `Float64` and remains as a `Float64` when it gets returned. Type Stability is important because if the Julia compiler can determine that the types of the variables stay constant for the lifetime of the function, it can perform optimizations on the machine code that would otherwise be impossible. These optimizations are evident when benchmarking the functions above. Using the `@benchmark` macro from `BenchmarkTools`, I timed both functions. They were run multiple times to ensure that only the compiled versions of the function were being executed rather than the functions being interpreted. The benchmarks showed that on average the type stable function ran in 306.178ns whereas the type unstable function ran in 573.355ns; almost 1.87x slower. When considering the many hundreds of function calls required to produce an output in `MERIT.jl`, these slowdowns become significant. When coding functions for `MERIT.jl`, I aimed to preserve type stability in functions that I knew would be called frequently. The `@code_warntype` macro was greatly beneficial in this task as it flags sections of code where the compiler cannot definitively infer the type of a variable or where it suspects there might be some type instability. Having type stable functions in `MERIT.jl` was a necessity, where possible, due to the large datasets that need to be processed. An 8cm radius breast at 0.25cm resolution has about 70k points, which need to have their distances to each antenna calculated. Assuming we are dealing with the MARIA M4 system with its 60 antennas, this equates to around 4.2M distance calculations alone for each image. Running such calculations on type unstable functions would be prohibitively expensive and would take far too long to compute. Herein lies one of the drawbacks of using the Julia language. While

the dynamism offers us powerful expressibility, it also comes at the cost of massive slowdowns when used incorrectly. While it may be easy to create type stable functions for the example shown in the code block above, it becomes increasingly harder to create type stable code when the operations that need to be performed become more complex. Creating performant code in these situations requires the developer to have advanced knowledge of the language. This does raise the barrier of entry for people who want to contribute to the library. However, I feel that this is an acceptable trade-off since the benefits that come from having a performant library, such as increased usage and better recognition, far outweigh the negatives from having an increased barrier of entry.

Bibliography

- [1] G. Stoet, PsyToolkit Testimonials, Std. [Online]. Available: https://www.psytoolkit.org/#_testimonials
- [2] A. W. Preece, I. Craddock, M. Shere, L. Jones, and H. L. Winton, "MARIA M4: Clinical evaluation of a prototype ultrawideband radar scanner for breast cancer detection." vol. 3, no. 3, p. 033502.
- [3] B. M. Moloney, P. F. McAnena, S. M. Elwahab, A. Fasoula, L. Duchesne, J. D. Gil Cano, C. Glynn, A. O'Connell, R. Ennis, A. J. Lowery, and M. J. Kerin, "The Wavelia Microwave Breast Imaging system-tumour discriminating features and their clinical usefulness." vol. 94, no. 1128, p. 20210907.
- [4] E. C. Fear, J. Bourqui, C. Curtis, D. Mew, B. Docktor, and C. Romano, "Microwave Breast Imaging With a Monostatic Radar-Based System: A Study of Application to Patients," vol. 61, no. 5, pp. 2119–2128.
- [5] M. Klemm, I. Craddock, J. Leendertz, A. Preece, and R. Benjamin, "Improved Delay-and-Sum Beamforming Algorithm for Breast Cancer Detection," vol. 2008.
- [6] B. Maklad, C. Curtis, E. Fear, and G. Messier, "Neighborhood-Based Algorithm to Facilitate the Reduction of Skin Reflections in Radar-Based Microwave Imaging," vol. 39, pp. 115–139.
- [7] J. Bourqui, J. Sill, and E. Fear, "A Prototype System for Measuring Microwave Frequency Reflections from the Breast," vol. 2012, p. 851234.